

# Unit 3

## Integrity Constraints and Relational Database Design

# Introduction

- ensure that **changes made to the database** by authorized users do not result in a loss of data consistency.
- integrity constraints guard against accidental damage to the database.
- seen two forms of integrity constraints for the E-R model
- **Key declarations** - certain attributes form a candidate key for a given entity set.
- **Form of a relationship** - many to many, one to many, one to one.

- integrity constraints with minimal overhead
- Domain constraints
- Referential Integrity
- Functional Dependency
- Normalization

# Domain Constraints

- **domain of possible values** must be associated with every attribute
- standard domain types, such as integer types, character types, and date/time types defined in SQL
- Declaring an attribute to be of a particular domain **acts as a constraint on the values** that it can take
- Domain constraints are the most **elementary form** of integrity constraint.
- They are tested easily by the **system whenever a new data item is entered** into the database

- for several attributes may have the same domain.
- For example, the attributes *customer-name* and *employee-name* might have the same domain: the set of all person names
- whether *customer-name* and *branch-name* should have the same domain.
- At the implementation level, both customer names and branch names are character strings.
- at the *conceptual*, rather than the physical level, *customer-name* and *branch-name* should have *distinct domains*

- **Equivalent Statement in Oracle**
- CREATE TABLE gender\_domain
  - (gender VARCHAR2(1) PRIMARY KEY,  
CONSTRAINT ch\_gen CHECK (gender IN ('M', 'F')));

# Referential Integrity

---

- value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
- This condition is called **referential integrity**.

# Referential Integrity in SQL

- Foreign keys can be specified as part of the SQL **create table** statement by using the **foreign key** clause

```
create table account  
( ...  
foreign key (branch-name) references branch(name)  
on delete cascade  
on update cascade,  
... )
```

- create table suppliers1(  
id number,name varchar2(50),CONSTRAINT FK\_SP1  
foreign key(id) references consumers(p\_id) on delete  
cascade )



- **on delete cascade**
- if a delete of a tuple in *branch* results in this referential-integrity constraint being violated, the system does not reject the delete.
- Instead, the delete “cascades” to the *account* relation, **deleting the tuple that refers to the branch that was deleted**
- Similarly, the system does not reject an update to field referenced by the constraint
- instead, the system updates the field **branch-name** in referencing tuples in *account* to new value

- actions other than **cascade**, if the constraint is violated:
- The referencing field (branch-name) can be set to null by using **set null in place of cascade**
- create table suppliers2  
(id number,name varchar2(50),  
CONSTRAINT FK\_SP2 foreign key(id) references  
consumers(p\_id) on delete set null )

# Database Modification

- Database modifications can cause violations of referential integrity.

$$\Pi_{\alpha}(r_2) \subseteq \Pi_K(r_1)$$

- Insert.** If a tuple  $t_2$  is inserted into  $r_2$ , the system must ensure that there is a tuple  $t_1$  in  $r_1$  such that  $t_1[K] = t_2[\alpha]$ . That is,

$$t_2[\alpha] \in \Pi_K(r_1)$$

- Delete.** If a tuple  $t_1$  is deleted from  $r_1$ , the system must compute the set of tuples in  $r_2$  that reference  $t_1$ :

$$\sigma_{\alpha = t_1[K]}(r_2)$$

# Database Modification

- **Update.** We must consider two cases for update: updates to the referencing relation ( $r_2$ ), and updates to the referenced relation ( $r_1$ ).
  - If a tuple  $t_2$  is updated in relation  $r_2$ , and the update modifies values for the foreign key  $\alpha$ , then a test similar to the insert case is made. Let  $t_2'$  denote the new value of tuple  $t_2$ . The system must ensure that

$$t_2'[\alpha] \in \Pi_K(r_1)$$

- If a tuple  $t_1$  is updated in  $r_1$ , and the update modifies values for the primary key ( $K$ ), then a test similar to the delete case is made. The system must compute

$$\sigma_{\alpha = t_1[K]}(r_2)$$

# Assertions

- An **assertion** is a predicate expressing a **condition that we wish the database always to satisfy**.
- Domain constraints and referential-integrity constraints are **special forms of assertions**.
- The sum of all loan amounts for each branch must be less than the sum of all account balances at the branch.
- Every loan has at least one customer who maintains an account with a **minimum balance of \$1000.00**.

- When an assertion is created, the **system tests it for validity.**
- If the assertion is valid, then any future modification to the database is allowed only if it **does not cause that assertion to be violated.**
- **This testing may introduce a significant amount of overhead if complex assertions have been made.**
- Hence, assertions should be used **with great care.**
- The high overhead of testing and maintaining assertions has led some system developers to **omit support for general assertions,**
- or to provide specialized forms of assertions that are easier to test

# Triggers

- A **trigger** is a statement that the **system executes automatically** as a side **effect of a modification** to the database.
- To design a trigger mechanism, we must meet two requirements:
  1. **Specify when a trigger is to be executed.**

This is broken up into an event that causes the trigger to be checked and a condition that must be satisfied for trigger execution to proceed.
  2. **Specify the *actions* to be taken when the trigger executes.**

- This model of triggers is referred to as the **event-condition-action model** for triggers.
- The database stores triggers just as if they were **regular data**, so that they are persistent and are accessible to all database operations.
- Once we enter a trigger into the database, the database system takes on the **responsibility of executing it** whenever the specified event occurs and the corresponding condition is satisfied.



# Need for Triggers

- Triggers are **useful mechanisms for alerting humans** or for starting certain tasks automatically when certain conditions are met.
- Example
- Instead of allowing **negative account balances**,
- the bank deals with overdrafts by setting the account **balance to zero**,
- and creating a loan in the amount of the overdraft.
- The bank gives this loan a loan number identical to the account number of the overdrawn account.

- For this example, the condition for executing the trigger is an **update to the *account* relation that results in a negative *balance* value.**
- Suppose that Jones' withdrawal of some money from an account made the account balance negative.
- Let  $t$  denote the account tuple with a negative *balance* value.
- The actions to be taken are:

- Insert a new tuple  $s$  in the loan relation with
$$s[\text{loan-number}] = t[\text{account-number}]$$
$$s[\text{branch-name}] = t[\text{branch-name}]$$
$$s[\text{amount}] = -t[\text{balance}]$$

(since  $t[\text{balance}]$  is negative, we negate  $t[\text{balance}]$  to get the loan amount—a positive number.)

- Insert a new tuple  $u$  in the borrower relation with
$$u[\text{customer-name}] = \text{"Jones"}$$
$$u[\text{loan-number}] = t[\text{account-number}]$$
- Set  $t[\text{balance}]$  to 0.

# Triggers in SQL

```
create trigger overdraft-trigger after update on account
referencing new row as nrow
for each row
when nrow.balance < 0
begin atomic
    insert into borrower
        (select customer-name, account-number
         from depositor
         where nrow.account-number = depositor.account-number);
    insert into loan values
        (nrow.account-number, nrow.branch-name, - nrow.balance);
    update account set balance = 0
        where account.account-number = nrow.account-number
end
```

# Authorization

---

- data stored in the database need protection from unauthorized access and malicious destruction or alteration
- Among the forms of malicious access are:
- Unauthorized reading of data (theft of information)
- Unauthorized modification of data
- Unauthorized destruction of data

- **Database security** refers to protection from malicious access
- To protect the database, we must take security measures at several levels:
- **Database system**
- **Operating system**
- **Network**
- **Physical**
- **Human**

# Authorization

---

- **Read authorization** allows reading, but not modification, of data.
- **Insert authorization** allows insertion of new data, but not modification of existing data.
- **Update authorization** allows modification, but not deletion, of data.
- **Delete authorization** allows deletion of data.

# Audit Trails

---

- An audit trail is a log of all changes (inserts/deletes/updates) to the database,
- along with information such as which user performed the change and when the change was performed.
- Many secure database applications require an **audit trail** be maintained.



- The audit trail aids security in several ways.
- For instance, if the balance on an account is found to be incorrect,
- the bank may wish to trace all the updates performed on the account, to find out incorrect (or fraudulent) updates, as well as the persons who carried out the updates.
- The bank could then also use the audit trail to trace all the updates performed by these persons, in order to find other incorrect or fraudulent updates

- It is possible to create an audit trail by defining appropriate triggers on relation updates
- many database systems provide built-in mechanisms to create audit trails,
- Which are much more convenient to use.
- Details of how to create audit trails vary across database systems,
- Can be found in database system manuals

# Authorization in SQL

---

- **Granting Privileges**
- **grant** <privilege list> **on** <relation name or view name> **to** <user/role list>
- **grant select on** *account* **to** *U1, U2, U3*
- **grant update** (*amount*) **on** *loan* **to** *U1, U2, U3*
- **grant references** (*branch-name*) **on** *branch* **to** *U1*

- **Roles**
- Roles can be created as follows
- **create role** *teller*
- Roles can then be granted privileges just as the users can:
- **grant select on** *account* **to** *teller*

**grant** *teller* **to** john  
**create role** *manager*  
**grant** *teller* **to** *manager*  
**grant** *manager* **to** mary

- **The Privilege to Grant Privileges**
- **grant select on *branch* to U1 with grant option**
- **revoke** <privilege list> **on** <relation name or view name> **from** <user/role list> [**restrict** / **cascade**]
- **Restrict -**  
system returns an error if there are any cascading revokes

# Pitfalls in Relational-Database Design

- Repetition of information
- Complicating to update
- Inability to represent certain information

*Lending-schema = (branch-name, branch-city, assets, customer-name, loan-number, amount)*

$t[\text{assets}]$  is the asset figure for the branch named  $t[\text{branch-name}]$ .

$t[\text{branch-city}]$  is the city in which the branch named  $t[\text{branch-name}]$  is located.

$t[\text{loan-number}]$  is the number assigned to a loan made by the branch named  $t[\text{branch-name}]$  to the customer named  $t[\text{customer-name}]$ .

$t[\text{amount}]$  is the amount of the loan whose number is  $t[\text{loan-number}]$ .

# Pitfalls in Relational-Database Design

- Add a new loan to our database
- loan is made by the Perryridge branch to Adams in the amount of \$1500. Let the loan-number be L-31
- Repeat the asset and city data for the Perryridge branch
- (Perryridge, Horseneck, 1700000, Adams, L-31, 1500)
- Asset of branch downtown changes from 17000 to 19000.
- Expected only one tuple need to change value.
- But In alternate design, more tuple gets changed & its costly

# Pitfalls in Relational-Database Design

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>	<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Downtown	Brooklyn	9000000	Jones	L-17	1000
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
Downtown	Brooklyn	9000000	Jackson	L-14	1500
Mianus	Horseneck	400000	Jones	L-93	500
Round Hill	Horseneck	8000000	Turner	L-11	900
Pownal	Bennington	300000	Williams	L-29	1200
North Town	Rye	3700000	Hayes	L-16	1300
Downtown	Brooklyn	9000000	Johnson	L-18	2000
Perryridge	Horseneck	1700000	Glenn	L-25	2500
Brighton	Brooklyn	7100000	Brooks	L-10	2200

Sample *lending* relation.

Prof. V. V. Kheradkar



# Pitfalls in Relational-Database Design

---

- Another problem with the Lending-schema design is that we cannot represent directly the information concerning a branch (branch-name, branch-city, assets) unless there exists at least one loan at the branch.

# Functional Dependencies

---

- the goal of a relational-database design is to generate a set of **relation schemas**
- that allows us to **store information without unnecessary redundancy**,
- yet also allows us to **retrieve information easily**.

One approach is to design schemas that are in an appropriate **normal form**.

- **Functional dependencies** play a key role in differentiating good database designs from bad database designs.
- A **functional dependency** is a type of **constraint** that is a **generalization of the notion of key**
- Functional dependencies are constraints on the set of legal relations

- **Definition:**
- A functional dependency occurs when one attribute in a relation uniquely determines another attribute.
- This can be written  $A \rightarrow B$
- which would be the same as stating "B is functionally dependent upon A."

- Let R be the relation, and
- let x and y be the arbitrary subset of the set of attributes of R.
- Then we say that Y is functionally dependent on x – in symbol.

$$X \rightarrow Y$$

- (Read x functionally determines y)
- If and only if each x value in R has associated with it precisely one y value in R
- In other words
- Whenever two tuples of R agree on their x value, they also agree on their Y value.

- In a table listing **employee** characteristics including **Social Security Number (SSN) and name**,
- it can be said that **name is functionally dependent upon SSN** (or  $SSN \rightarrow name$ )
- because an employee's name can be uniquely determined from their SSN.
- However, the reverse statement ( $name \rightarrow SSN$ ) is not true because more than one employee can have the same name but different SSNs.

# Basic Concepts

- Consider a relation schema  $R$ , and
- let  $\alpha \subseteq R$  and  $\beta \subseteq R$ .
- The **functional dependency**  $\alpha \rightarrow \beta$  holds on schema  $R$  if,
- in any legal relation  $r(R)$ ,
- for all pairs of tuples  $t_1$  and  $t_2$  in  $r$  such that
- $t_1[\alpha] = t_2[\alpha]$ , then
- $t_1[\beta] = t_2[\beta]$ .

1:1 relationship between attribute(s) on left and right-hand side of a dependency; hold for all time;

- **$X \rightarrow Y$  means**
- Given any two tuples in  $r$ , if the  $X$  values are the same,
- then the  $Y$  values must also be the same.
- (but not vice versa)

• Read “ $\rightarrow$ ” as “determines”

if “ $K \rightarrow$  all attributes of  $R$ ”

then  $K$  is a superkey for  $R$

- **FDs are a generalization of keys.**



- Loan-info-schema = (loan-number, branch-name, customer-name, amount)
- The set of functional dependencies that we expect to hold on this relation schema is
- $\text{loan-number} \rightarrow \text{amount}$
- $\text{loan-number} \rightarrow \text{branch-name}$
- We would not, however, expect the functional dependency
- $\text{loan-number} \rightarrow \text{customer-name}$
- in general, a given loan can be made to more than one customer (for example, to both members of a husband–wife pair)

- We shall use functional dependencies in two ways:
- 1. To test relations to see whether they are legal under a given set of functional dependencies.
- If a relation  $r$  is legal under a set  $F$  of functional dependencies, we say that  $r$  **satisfies**  $F$ .

- **2. To specify constraints on the set of legal relations.**
- concern with only those relations that satisfy a given set of functional dependencies.
- If we wish to constrain ourselves to relations on schema  $R$  that satisfy a set  $F$  of functional dependencies,
- we say that  $F$  **holds** on  $R$ .

- Some functional dependencies are said to be **trivial**
- because they are satisfied by all relations.
- For example,  $A \rightarrow A$  is satisfied by all relations involving attribute  $A$ .
- for all tuples  $t_1$  and  $t_2$  such that  $t_1[A] = t_2[A]$ ,
- it is the case that  $t_1[A] = t_2[A]$ .
- Similarly,  $AB \rightarrow A$  is satisfied by all relations involving attribute  $A$ .
- In general, a functional dependency of the form  $\alpha \rightarrow \beta$  is **trivial** if  $\beta \subseteq \alpha$ .

- An FD is trivial if and only if
  - the right hand side is a subset of the left hand side.
  - e.g.  $\langle S\#, P\# \rangle \rightarrow \langle S\# \rangle$ . (Trivial)
  - **Nontrivial** dependencies are the one, which are not trivial.
- 
- when we design a relational database,
  - we first list those functional dependencies that must always hold.
  - In the banking example, list of dependencies includes the following:

- Branch-schema = (branch-name, branch-city, assets)
- Customer-schema = (customer-name, customer-street, customer-city)
- Loan-schema = (loan-number, branch-name, amount)
- Borrower-schema = (customer-name, loan-number)
- Account-schema = (account-number, branch-name, balance)
- Depositor-schema = (customer-name, account-number)

- On Branch-schema:
  - branch-name  $\rightarrow$  branch-city
  - branch-name  $\rightarrow$  assets
- On Customer-schema:
  - customer-name  $\rightarrow$  customer-city
  - customer-name  $\rightarrow$  customer-street
- On Loan-schema:
  - loan-number  $\rightarrow$  amount
  - loan-number  $\rightarrow$  branch-name

- On Borrower-schema:
  - No functional dependencies
- On Account-schema:
  - account-number  $\rightarrow$  branch-name
  - account-number  $\rightarrow$  balance
- On Depositor-schema:
  - No functional dependencies



# Types of functional dependencies:

---


- Full Functional dependency
- Partial Functional dependency
- Transitive dependency

# Full dependencies

*X and Y are attributes.*

*X Functionally determines Y*

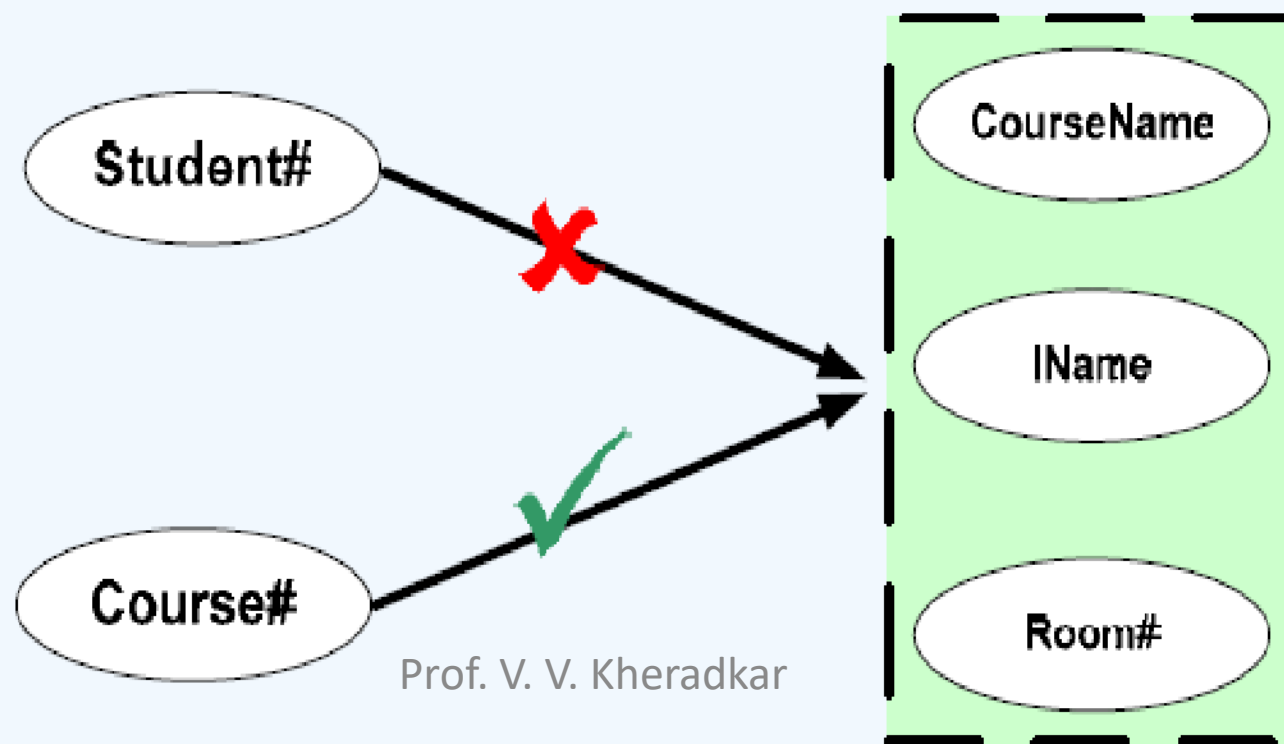
*Note: Subset of X should not functionally determine Y*



# Partial dependencies

*X and Y are attributes.*

*Attribute Y is partially dependent on the attribute X only if it is dependent on a sub-set of attribute X.*




# Transitive dependencies

*X Y and Z are three attributes.*

**$X \rightarrow Y$**

**$Y \rightarrow Z$**

**$\Rightarrow X \rightarrow Z$**



# Closure of a Set of Functional Dependencies

- The set of all **FDs** that are implied by a given set **S** of **FDs**
- is called the closure of S, denoted by **S<sup>+</sup>**
- It is not sufficient to consider the given set of functional dependencies.
- We need to consider **all functional dependencies** that hold

- given a relation schema  $R = (A, B, C, G, H, I)$
- and the set of functional dependencies
- $A \rightarrow B$
- $A \rightarrow C$
- $CG \rightarrow H$
- $CG \rightarrow I$
- $B \rightarrow H$
- The functional dependency  $A \rightarrow H$  is logically implied.

- Suppose
- $t_1$  and  $t_2$  are tuples such that
- $t_1[A] = t_2[A]$
- Since we are given that  $A \rightarrow B$ ,
- it follows from the definition of functional dependency that
- $t_1[B] = t_2[B]$
- Then, since we are given that  $B \rightarrow H$ ,
- it follows from the definition of functional dependency that
- $t_1[H] = t_2[H]$

- Let  $F$  be a set of functional dependencies.
- The **closure** of  $F$ , denoted by  $F^+$ ,
- is the set of **all functional dependencies logically implied by  $F$** .
- Given  $F$ , we can compute  $F^+$  directly from the formal definition of functional dependency.
- If  $F$  is large, this process would be **lengthy and difficult**.



- **Axioms, or rules of inference**, provide a simpler technique for reasoning about functional dependencies.
- We can use the following three rules to find **logically implied functional dependencies**.
- By applying these rules repeatedly, we can find all of  $F_+$ , given  $F$ .
- This collection of rules is called **Armstrong's axioms** in honor of the person who first proposed it.

- **Reflexivity rule.**
- If  $\alpha$  is a set of attributes and  $\beta \subseteq \alpha$ , then  $\alpha \rightarrow \beta$  holds
- **Augmentation rule.**
- If  $\alpha \rightarrow \beta$  holds and  $\gamma$  is a set of attributes, then  $\gamma\alpha \rightarrow \gamma\beta$  holds.
- **Transitivity rule.**
- If  $\alpha \rightarrow \beta$  holds and  $\beta \rightarrow \gamma$  holds, then  $\alpha \rightarrow \gamma$  holds.

- Armstrong's axioms are **sound**,
- because they **do not generate any incorrect functional dependencies**.
- They are **complete**, because, for a given set  $F$  of functional dependencies, they allow us to generate all  $F^+$ .

- Although Armstrong's axioms are complete,
- it is tiresome to use them directly for the computation of  $F_+$ .
- To simplify matters further, we list **additional rules**.

- **Union rule.**
- If  $\alpha \rightarrow \beta$  holds and  $\alpha \rightarrow \gamma$  holds, then  $\alpha \rightarrow \beta\gamma$  holds.
- **Decomposition rule.**
- If  $\alpha \rightarrow \beta\gamma$  holds, then  $\alpha \rightarrow \beta$  holds and  $\alpha \rightarrow \gamma$  holds.
- **Pseudotransitivity rule.**
- If  $\alpha \rightarrow \beta$  holds and  $\gamma\beta \rightarrow \delta$  holds, then  $\alpha\gamma \rightarrow \delta$  holds.

1. **Reflexivity:** if  $B$  is a subset of  $A$ , then  $A \rightarrow B$ .
2. **Augmentation:** if  $A \rightarrow B$  then  $AC \rightarrow BC$
3. **Transitivity:** if  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$ .
4. **Self – determination:**  $A \rightarrow A$ .
5. **Decomposition:** If  $A \rightarrow BC$ , then  $A \rightarrow B, A \rightarrow C$ .
6. **Union:** if  $A \rightarrow B$  and  $A \rightarrow C$ , then  $A \rightarrow BC$
7. **Composition:** if  $A \rightarrow B, C \rightarrow D$  then  $AC \rightarrow BD$ .
8. If  $A \rightarrow B$  and  $C \rightarrow D$ , then  $All (C - B) \rightarrow BD$

- Let us apply our rules to the example of schema
- $R = (A, B, C, G, H, I)$  and
- the set  $F$  of functional dependencies
- $\{A \rightarrow B, A \rightarrow C, CG \rightarrow H, CG \rightarrow I, B \rightarrow H\}$ .
- We list several members of  $F^+$  here:
- $A \rightarrow H$ . Since  $A \rightarrow B$  and  $B \rightarrow H$  hold, we apply the transitivity rule.
- $CG \rightarrow HI$ . Since  $CG \rightarrow H$  and  $CG \rightarrow I$ , the union rule implies that  $CG \rightarrow HI$ .

- $AG \rightarrow I$ . Since  $A \rightarrow C$  and  $CG \rightarrow I$ , the **pseudotransitivity** rule implies that  $AG \rightarrow I$  holds.
- Another way of finding that  $AG \rightarrow I$  holds is as follows. We use the augmentation rule on  $A \rightarrow C$  to infer  $AG \rightarrow CG$ .
- Applying the transitivity rule to this dependency and  $CG \rightarrow I$ , we infer  $AG \rightarrow I$ .



$F^+ = F$

repeat

  for each functional dependency  $f$  in  $F^+$

    apply reflexivity and augmentation rules on  $f$

    add the resulting functional dependencies to  $F^+$

  for each pair of functional dependencies  $f_1$  and  $f_2$  in  $F^+$

    if  $f_1$  and  $f_2$  can be combined using transitivity

      add the resulting functional dependency to  $F^+$

until  $F^+$  does not change any further

**Figure 7.6** A procedure to compute  $F^+$ .

# Closure of Attribute Sets

- To test whether a **set  $\alpha$  is a superkey**,
- we must devise an algorithm for computing the set of attributes functionally determined by  $\alpha$ .
- One way of doing this is to compute  $F_+$ ,
- take all functional dependencies with  **$\alpha$  as the left-hand side**, and take the **union** of the right-hand sides of all such dependencies.
- However, doing so can be **expensive**, since  $F_+$  can be large.

# Closure of Attribute Sets

- Algorithm to compute  $\alpha^+$ , the closure of  $\alpha$  under  $F$

*result* :=  $\alpha$ ;

**while** (changes to *result*) **do**

**for each** functional dependency  $\beta \rightarrow \gamma$  **in**  $F$  **do**

**begin**

**if**  $\beta \subseteq \text{result}$  **then** *result* := *result*  $\cup \gamma$ ;

**end**

- compute  $(AG)^+$  with the functional dependencies  $\{A \rightarrow B, A \rightarrow C, CG \rightarrow H, CG \rightarrow I, B \rightarrow H\}$ .
- We start with  $result = AG$ .
- $A \rightarrow B$  causes us to include  $B$  in result.
- we observe that  $A \rightarrow B$  is in  $F$ ,  $A \subseteq result$  (which is  $AG$ ), so  $result := result \cup B$ .
- $A \rightarrow C$  causes result to become  $ABCG$ .
- $CG \rightarrow H$  causes result to become  $ABCGH$ .
- $CG \rightarrow I$  causes result to become  $ABCGHI$ .

# Canonical Cover

- Suppose that we have a **set of functional dependencies  $F$**  on a relation schema.
- Whenever a user performs an update on the relation,
- the database system must ensure that the **update does not violate any functional dependencies**,
- that is, **all the functional dependencies in  $F$  are satisfied in the new database state**.
- The system must **roll back** the update if it violates any functional dependencies in the set  $F$ .

- We can **reduce the effort** spent in **checking for violations** by testing a **simplified set of functional dependencies** that has the **same closure as the given set**.
- Any database that satisfies the simplified set of functional dependencies will also satisfy the original set, and vice versa,
- since the two sets have the same closure.
- However, the simplified set is **easier to test**.

- An attribute of a functional dependency is said to be **extraneous**
- if we can remove it **without changing the closure** of the set of functional dependencies.
- The formal definition of **extraneous attributes** is as follows.
- For example, suppose we have the functional dependencies  $AB \rightarrow C$  and  $A \rightarrow C$  in  $F$ .
- Then,  $B$  is extraneous in  $AB \rightarrow C$ .

- A **canonical cover**  $F_c$  for  $F$  is a set of dependencies such that  $F$  logically implies all dependencies in  $F_c$ ,
- and  $F_c$  logically implies all dependencies in  $F$ .
- Furthermore,  $F_c$  must have the following properties:
- No functional dependency in  $F_c$  contains an extraneous attribute.
- Each left side of a functional dependency in  $F_c$  is unique.
- That is, there are no two dependencies  $\alpha_1 \rightarrow \beta_1$  and  $\alpha_2 \rightarrow \beta_2$  in  $F_c$  such that  $\alpha_1 = \alpha_2$ .



$F_c = F$

repeat

Use the union rule to replace any dependencies in  $F_c$  of the form

$\alpha_1 \rightarrow \beta_1$  and  $\alpha_1 \rightarrow \beta_2$  with  $\alpha_1 \rightarrow \beta_1 \beta_2$ .

Find a functional dependency  $\alpha \rightarrow \beta$  in  $F_c$  with an extraneous attribute either in  $\alpha$  or in  $\beta$ .

/\* Note: the test for extraneous attributes is done using  $F_c$ , not  $F^*$  \*/

If an extraneous attribute is found, delete it from  $\alpha \rightarrow \beta$ .

until  $F_c$  does not change.

**Figure 7.8** Computing canonical cover

- Consider the following set  $F$  of functional dependencies on schema  $(A,B,C)$ :
- $A \rightarrow BC$
- $B \rightarrow C$
- $A \rightarrow B$
- $AB \rightarrow C$
- Let us compute the canonical cover for  $F$ .

- There are two functional dependencies with the same set of attributes on the left side of the arrow:
- $A \rightarrow BC$
- $A \rightarrow B$
- We combine these functional dependencies into  $A \rightarrow BC$ .
- A is extraneous in  $AB \rightarrow C$  because F logically implies  $(F - \{AB \rightarrow C\}) \cup \{B \rightarrow C\}$ .
- This assertion is true because  $B \rightarrow C$  is already in our set of functional dependencies.

- C is extraneous in  $A \rightarrow BC$ , since  $A \rightarrow BC$  is logically implied by  $A \rightarrow B$  and  $B \rightarrow C$ .
- Thus, our canonical cover is
  - $A \rightarrow B$
  - $B \rightarrow C$
- A canonical cover might not be unique.

# Testing if an Attribute is Extraneous

Consider a set  $F$  of functional dependencies and the functional dependency  $\alpha \rightarrow \beta$  in  $F$ .

To test if attribute  $A \in \alpha$  is extraneous in  $\alpha$

1. compute  $(\{\alpha\} - A)^+$  using the dependencies in  $F$
2. check that  $(\{\alpha\} - A)^+$  contains  $\beta$ ; if it does,  $A$  is extraneous in  $\alpha$

To test if attribute  $A \in \beta$  is extraneous in  $\beta$

1. compute  $\alpha^+$  using only the dependencies in  $F' = (F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$ ,
2. check that  $\alpha^+$  contains  $A$ ; if it does,  $A$  is extraneous in  $\beta$

# Computing a Canonical Cover

*Q.1. Let  $R = (A, B, C)$  and  $F = \{AB \rightarrow C, A \rightarrow C\}$   
Find canonical cover?*

*Check  $A$  is extraneous in  $AB \rightarrow C$ ?  
compute  $(AB - A)^+$  under  $F$   
if contains  $C$  then extraneous  
Check  $B$  is extraneous in  $AB \rightarrow C$ ?  
compute  $(AB - B)^+$  under  $F$   
if contains  $C$  then extraneous*

*Q.2. Let  $R = (A, B, C, D, E)$  and  $F = \{AB \rightarrow CD, A \rightarrow E, E \rightarrow C\}$   
Find canonical cover?*

*Check  $A$  is extraneous in  $AB \rightarrow CD$ ?  
Check  $B$  is extraneous in  $AB \rightarrow CD$ ?*

# Decomposition

- Lending-schema = (branch-name, branch-city, assets, customer-name, loan-number, amount)
- The set  $F$  of functional dependencies that we require to hold on Lending-schema are
  - branch-name  $\rightarrow$  branch-city, assets
  - loan-number  $\rightarrow$  amount, branch-name

- Lending-schema is an example of a bad database design.
- Assume that we decompose it to the following three relations:
- Branch-schema = (branch-name, branch-city, assets)
- Loan-schema = (loan-number, branch-name, amount)
- Borrower-schema = (customer-name, loan-number)



# Normalization


---

- the goal of a relational-database design is to generate a set of relation schemas
- that allows us to store information without unnecessary redundancy,
- yet also allows us to retrieve information easily.
- One approach is to design schemas that are in an appropriate normal form.
- design standards are called as normal forms

- Database design theory includes design standards called *normal forms*.
- The process of making your data and tables match these standards is called *normalizing data or data normalization*.
- By normalizing your data, you **eliminate redundant information**
- and **organize your table to make it easier to manage the data and make future changes to the table and database structure.**
- This process removes the insertion, deletion, and modification anomalies.

- In normalizing data, we usually **divide large tables into smaller, easier to maintain tables.**
- we can then use the technique of **adding foreign keys to enable connections between the tables.**
- Data normalization is part of the **database design process** and is not specific nor unique to any particular RDBMS.
- There are first, second, third, Boyce-Codd, fourth, and fifth normal forms.
- Each normal form represents an increasingly stringent set of rules;

# Relationship Between Normal Forms



# Unnormalized Form (UNF)

---

- A table that contains one or more repeating groups.
- To create an unnormalized table:
- transform data from information source (e.g. form) into table format with columns and rows.

# First Normal Form

- **First normal form**, imposes a very basic requirement on relations;
- unlike the other normal forms, it does not require additional information such as functional dependencies.
- **A relation in which intersection of each row and column contains one and only one value.**
- A table is in first normal form (1NF) if there are **no repeating groups**.
- A repeating group is a set of logically related fields or values that occur multiple times in one record.

- Under first normal form, all occurrences of a record type must contain the **same number of fields**.
- For all attributes only one value
- All the attributes are atomic.
- **Atomic**- the smallest level to which data may be broken down and remain meaningful.

- 
- example
- 1. if the schema of a relation **employee** included an attribute **children** whose domain elements are sets of names,
- the schema would not be in first normal form.
- 2. Composite attributes, such as an attribute **address** with component attributes street and city, also have nonatomic domains.



EmployeeID	Name	Project	Time
EN1-26	Sean O'Brien	30-452-T3, 30-457-T3, 32-244-T3	0.25, 0.40, 0.30
EN1-33	Amy Guya	30-452-T3, 30-382-TC, 32-244-T3	0.05, 0.35, 0.60
EN1-35	Steven Baranco	30-452-T3, 31-238-TC	0.15, 0.80
EN1-36	Elizabeth Roslyn	35-152-TC	0.90
EN1-38	Carol Schaaf	36-272-TC	0.75
EN1-40	Alexandra Wing	31-238-TC, 31-241-TC	0.20, 0.70

EmpID	Last Name	First Name	Project1	Time1	Project2	Time2	Project3	Time3
EN1-26	O'Brien	Sean	30-452-T3	0.25	30-457-T3	0.40	32-244-T3	0.30
EN1-33	Guya	Amy	30-452-T3	0.05	30-382-TC	0.35	32-244-T3	0.60
EN1-35	Baranco	Steven	30-452-T3	0.15	31-238-TC	0.80		
EN1-36	Roslyn	Elizabeth	35-152-TC	0.90				
EN1-38	Schaaf	Carol	36-272-TC	0.75				
EN1-40	Wing	Alexandra	31-238-TC	0.20	31-241-TC	0.70		

* EmployeeID	Last Name	First Name
EN1-26	O'Brien	Sean
EN1-33	Guya	Amy
EN1-35	Baranco	Steven
EN1-36	Roslyn	Elizabeth
EN1-38	Schaaf	Carol
EN1-40	Wing	Alexandra

* ProjectNum	* EmployeeID	Time
30-328-TC	EN1-33	0.35
30-452-T3	EN1-26	0.25
30-452-T3	EN1-33	0.05
30-452-T3	EN1-35	0.15
31-238-TC	EN1-35	0.80
30-457-T3	EN1-26	0.40
31-238-TC	EN1-40	0.20
31-241-TC	EN1-40	0.70
32-244-T3	EN1-33	0.60
35-152-TC	EN1-36	0.90
36-272-TC	EN1-38	0.75

# Second Normal Form (2NF)

---

- Suppose a table is in first normal form and
- each non-key field is functionally dependent on the entire primary key.
- Look for values that occur multiple times in a non-key field.
- This tells you that you have too many fields in a single table.
- No Partial dependency between non-key attributes and key attributes

*EmployeeID	LastName	FirstName	*ProjectNumber	ProjectTitle
EN1-26	O'Brien	Sean	30-452-T3	STAR manual
EN1-26	O'Brien	Sean	30-457-T3	ISO procedures
EN1-26	O'Brien	Sean	31-124-T3	Employee handbook
EN1-33	Guya	Amy	30-452-T3	STAR manual
EN1-33	Guya	Amy	30-482-TC	Web Site
EN1-33	Guya	Amy	31-241-TC	New catalog
EN1-35	Baranco	Steven	30-452-T3	STAR manual
EN1-35	Baranco	Steven	31-238-TC	STAR prototype
EN1-36	Roslyn	Elizabeth	35-152-TC	STAR pricing
EN1-38	Schaaf	Carol	36-272-TC	Order system
EN1-40	Wing	Alexandra	31-238-TC	STAR prototype
EN1-40	Wing	Alexandra	31-241-TC	New catalog

- Solution to this lies in breaking the table into smaller tables.
- more tables is the solution to most problems encountered during data normalisation.
- This removes the modification anomaly of having the repeated values.

## **Complying with second normal form**

design new tables that will eliminate the repeated data in the non-key fields.

- 1. Decide what fields belong together in a table, think about which field determines the values in other fields.
- Create a table for those fields and enter the sample data.
- 2. Think about what the primary key for each table would be and about the relationship between the tables.
- If necessary, add foreign keys.
- 3. Mark the primary key for each table and make sure that you don't have repeating data in non-key fields.



# EMPLOYEES

*EmployeeID	Last Name	First Name
EN1-26	O'Brien	Sean
EN1-33	Guya	Amy
EN1-35	Baranco	Steven
EN1-36	Roslyn	Elizabeth
EN1-38	Schaaf	Carol
EN1-40	Wing	Alexandra

- Employee\_Projects

*EmployeeID	*ProjectNum
EN1-26	30-452-T3
EN1-26	30-457-T3
EN1-26	31-124-T3
EN1-33	30-328-TC
EN1-33	30-452-T3
EN1-33	32-244-T3
EN1-35	30-452-T3
EN1-35	31-238-TC
EN1-36	35-152-TC
EN1-38	36-272-TC
EN1-40	31-238-TC
EN1-40	31-241-TC

- Projects

*ProjectNum	ProjectTitle
30-452-T3	STAR manual
30-457-T3	ISO procedures
30-482-TC	Web site
31-124-T3	Employee handbook
31-238-TC	STAR prototype
31-238-TC	New catalog
35-152-TC	STAR pricing
36-272-TC	Order system

# Third Normal Form (3NF)

- Consider a table is in second normal form (2NF) and there are **no transitive dependencies**.
- A **transitive dependency** is a type of functional dependency in which
- the value in a **non-key field** is determined by the value in **another non-key field** and that field is not a candidate key.
- **3NF - A relation that is in 1NF and 2NF and in which no non-key attribute is transitively dependent on the primary key.**

*ProjectNum	ProjectTitle	ProjectMgr	Phone
30-452-T3	STAR manual	Garrison	2756
30-457-T3	ISO procedures	Jacanda	2954
30-482-TC	Web site	Friedman	2846
31-124-T3	Employee handbook	Jones	3102
31-238-TC	STAR prototype	Garrison	2756
31-241-TC	New catalog	Jones	3102
35-152-TC	STAR pricing	Vance	3022
36-272-TC	Order system	Jacanda	2954

- The phone number is repeated each time a manager name is repeated.
- This is because the phone number is only a second cousin to the project number.
- It's dependent on the manager,
- which is dependent on the project number
- (a transitive dependency).

- The ProjectMgr field is not a **candidate key** because
- the same person manages more than one project.
- Again, the solution is to remove the field with repeating data to a separate table.

- **Complying with third normal form**
- 1. Think about which fields belong together and create new tables to hold them.
- 2. Enter the sample data and check for unnecessarily (not part of primary key) repeated values.
- 3. Identify the primary key for each table and, if necessary, add foreign keys.



- Projects

*ProjectNum	ProjectTitle	ProjectMgr
30-452-T3	STAR manual	Garrison
30-457-T3	ISO procedures	Jacanda
30-482-TC	Web site	Friedman
31-124-T3	Employee handbook	Jones
31-238-TC	STAR prototype	Garrison
31-241-TC	New catalog	Jones
35-152-TC	STAR pricing	Vance
36-272-TC	Order system	Jacanda

- # Manager

*ProjectMgr	Phone
Friedman	2846
Garrison	2756
Jacanda	2954
Jones	3102
Vance	3022

- there are no unnecessarily repeating values in non-key fields
- and the value in each non-key field is determined by the value(s) in the key field(s).

# Boyce-Codd Normal Form

- A table is in third normal form (3NF) and all determinants are candidate keys.
- Boyce-Codd normal form (BCNF) can be thought of as a "new" third normal form.
- It was introduced to cover situations that the "old" third normal form did not address.
- determinant - determines the value in another field
- candidate keys -qualify for designation as primary key.
- This normal form applies to situations where you have overlapping candidate keys.

# Candidate Key

---

- A **superkey** is a set of one or more attributes that, taken collectively, allow us to identify uniquely an entity in the entity set.
- superkeys for which no proper subset is a superkey.
- Such minimal superkeys are called **candidate keys**.
- several distinct sets of attributes could serve as a candidate key.
- **primary key** denote a candidate key that is chosen by the database designer as the principal means of identifying entities within an entity set.

Student#	EmailID	Course#	Marks
101	<a href="mailto:Davis@myuni.edu">Davis@myuni.edu</a>	M4	82
102	<a href="mailto:Daniel@myuni.edu">Daniel@myuni.edu</a>	M4	62
101	<a href="mailto:Davis@myuni.edu">Davis@myuni.edu</a>	H6	79
103	<a href="mailto:Sandra@myuni.edu">Sandra@myuni.edu</a>	C3	65
104	<a href="mailto:Evelyn@myuni.edu">Evelyn@myuni.edu</a>	B3	77
102	<a href="mailto:Daniel@myuni.edu">Daniel@myuni.edu</a>	P3	68
105	<a href="mailto:Susan@myuni.edu">Susan@myuni.edu</a>	P3	89
103	<a href="mailto:Sandra@myuni.edu">Sandra@myuni.edu</a>	B4	54
105	<a href="mailto:Susan@myuni.edu">Susan@myuni.edu</a>	H6	87
104	<a href="mailto:Evelyn@myuni.edu">Evelyn@myuni.edu</a>	M4	65

- In the above table, we have two candidate keys namely
- **STUDENT# COURSE#**
- **COURSE# EmailId.**
- COURSE# is overlapping among those candidate keys.
- Hence these candidate keys are called as **“Overlapping Candidate Keys”**.

- there are four determinants in this relation namely:
  - STUDENT# (STUDENT# decides EmailID)
  - EmailID (EmailID decides STUDENT#)
  - STUDENT# COURSE# (decides Marks)
  - COURSE# EmailID (decides Marks).
- 
- Only combination of STUDENT# COURSE# and COURSE# EmailID are candidate keys.



## STUDENT TABLE

Student#	EmailID
101	<a href="mailto:Davis@myuni.edu">Davis@myuni.edu</a>
102	<a href="mailto:Daniel@myuni.edu">Daniel@myuni.edu</a>
103	<a href="mailto:Sandra@myuni.edu">Sandra@myuni.edu</a>
104	<a href="mailto:Evelyn@myuni.edu">Evelyn@myuni.edu</a>
105	<a href="mailto:Susan@myuni.edu">Susan@myuni.edu</a>

<b>Student#</b>	<b>Course#</b>	<b>Marks</b>
<b>101</b>	<b>M4</b>	<b>82</b>
<b>102</b>	<b>M4</b>	<b>62</b>
<b>101</b>	<b>H6</b>	<b>79</b>
<b>103</b>	<b>C3</b>	<b>65</b>
<b>104</b>	<b>B3</b>	<b>77</b>
<b>102</b>	<b>P3</b>	<b>68</b>
<b>105</b>	<b>P3</b>	<b>89</b>
<b>103</b>	<b>B4</b>	<b>54</b>
<b>105</b>	<b>H6</b>	<b>87</b>
<b>104</b>	<b>M4</b>	<b>65</b>

- Difference between 3NF and BCNF is that for a functional dependency  $A \rightarrow B$ ,
- BCNF insists that for this dependency to remain in a relation,  $A$  must be a candidate key.
- Every relation in BCNF is also in 3NF. However, relation in 3NF may not be in BCNF.

# Fourth Normal Form:

---

- A table is in Boyce-Codd normal form (BCNF) and there are no multi-valued dependencies.
- A multi-valued dependency occurs when,
- for each value in field A, there is a set of values for field B and a set of values for field C but fields B and C are not related.

- Although BCNF removes anomalies due to functional dependencies,
- another type of dependency called a **multi-valued dependency (MVD)** can also cause data redundancy.
- Possible existence of MVDs in a relation is due to 1NF and can result in data redundancy.

- Look for repeated or null values in non-key fields.
- A multi-valued dependency occurs when the table contains fields that are not logically related.

*Movie	*Star	*Producer
Once Upon a Time	Julie Garland	Alfred Brown
Once Upon a Time	Mickey Rooney	Alfred Brown
Once Upon a Time	Julie Garland	Muriel Humphreys
Once Upon a Time	Mickey Rooney	Muriel Humphreys
Moonlight	Humphrey Bogart	Alfred Brown
Moonlight	Julie Garland	Alfred Brown

- A movie can have **more than one star and more than one producer.**
- A star can be in more than one movie.
- A producer can produce more than one movie.
- The primary key would have to include all three fields and so this table would be in **BCNF.**
- have unnecessarily repeated values, with the data maintenance problems that causes,
- and will have trouble with deletion anomalies.

- The Star and the Producer really aren't logically related.
- The Movie determines the Star and the Movie determines the Producer.
- The answer is to have a separate table for each of those logical relationships - **one holding Movie and Star and the other with Movie and Producer**, as shown below:



\*Movie

Once Upon a Time

Once Upon a Time

Moonlight

Moonlight

\*Movie

Once Upon a Time

Once Upon a Time

Moonlight

\*Star

Julie Garland

Mickey Rooney

Humphrey Bogart

Julie Garland

\*Producer

Alfred Brown

Muriel Humphreys

Alfred Brown

# Fifth Normal Form

- A table is in fourth normal form (4NF) and there are **no cyclic dependencies**.
- A cyclic dependency can occur only when you have a **multi-field primary key consisting of three or more fields**.
- For example, let's say your primary key consists of fields A, B, and C.
- A cyclic dependency would arise if the values in those fields were related in pairs of A and B, B and C, and A and C.

- Fifth normal form is also called **projection-join** normal form.
- A **projection** is a new table holding a subset of fields from an original table.
- When properly formed projections are **joined**, they must **result in the same set of data** that was contained in the original table

- Look for the number of records that will have to be added or maintained
- Following is some sample data about buyers, the products they buy, and the companies they buy from on behalf of MegaMall, a large department store.

Buyer	*Product	*Company
Chris	jeans	Levi
Chris	jeans	Wrangler
Chris	shirts	Levi
Lori	jeans	Levi

- a table with cyclic dependencies
- The primary key consists of all three fields.
- One data maintenance problem that occurs is that you need to **add a record for every buyer who buys a product for every company that makes that product** or they can't buy from them.
- That may not appear to be a big deal in this sample of 2 buyers, 2 products, and 2 companies ( $2 \times 2 \times 2 = 8$  total records).
- But what if you went to 20 buyers, 50 products, and 100 companies ( $20 \times 50 \times 100 = 100,000$  potential records)?

\*Buyer

Chris

Chris

Lori

Product

jeans

jeans

shirts

\*Product

jeans

shirts

jeans

\*Company

Wrangler

Levi

Levi

- However, if you joined the two tables above on the Product field,
- it would produce a record not part of the original data set
- it would say that Lori buys jeans from Wrangler.
- This is where the **projection-join** concept comes in.

\*Buyer

Chris

Chris

Lori

\*Product

jeans

shirts

jeans

\*Product

jeans

jeans

shirts

\*Company

Wrangler

Levi

Levi

\*Buyer

Chris

Chris

Lori

\*Company

Levi

Wrangler

Levi



- Above, tables that comply with 5NF
- When the first two tables are joined by Product and the result joined to the third table by Buyer and Company, the result is the original set of data.
- In our scenario of 20 buyers, 50 products, and 100 companies, you would have, at most, 1000 records in the Buyers table ( $20 \times 50$ ), 5000 records in the Products table ( $50 \times 100$ ), and 2000 records in the Companies table ( $20 \times 100$ ). With a maximum of 8000 records, these tables would be much easier to maintain than the possible 100,000 records we saw earlier.

# Denormalization

---

- Normalization is one of many database design goals
- Normalized table requirements
  - Additional processing
  - Loss of system speed

Normalization purity is difficult to sustain due to conflict in:

Design efficiency  
Information requirements  
Processing

# Unnormalized Table Defects

---

- Data updates less efficient
- Indexing more cumbersome
- No simple strategies for creating views